Lab4
Neural Network

Chia-Suan Yu
Po-Chih Kuo

Introduction

In this lab, you will apply neural networks to three tasks:

- **Regression:** Approximate a mathematical function.
- **Binary Classification:** Classify OCT retinal images as either normal or affected by Choroidal Neovascularization (CNV), a condition that can cause vision loss.
- **Multi-Class Classification:** Classify OCT retinal images into four categories:
  - CNV: Abnormal blood vessel growth in the eye
  - DME: Swelling of the macula due to diabetes
  - Drusen: Yellow deposits under the retina, potential sign of age-related macular degeneration
  - Normal: Healthy retina

Goal

- Build your own deep neural network step by step
- Implement all the functions required to build a deep neural network
- Understanding forward propagation, backward propagation and update
- Implement Binary Cross-Entropy loss and Categorical Cross-Entropy loss
- Implement **regression** (basic part), **binary classifier** (basic part) and **multi-class classifier** (advanced part)

Grading Policy

Item	Score
Basic Implementation	65%
Advanced Implementation	30%
Basic & Advanced Report	5%

Overview

Layer with parameters

Dense layer

- initialize_parameters
- forward
- backward
- update

Layer without parameters

Activation function layer

- forward
- backward

Overview

Model

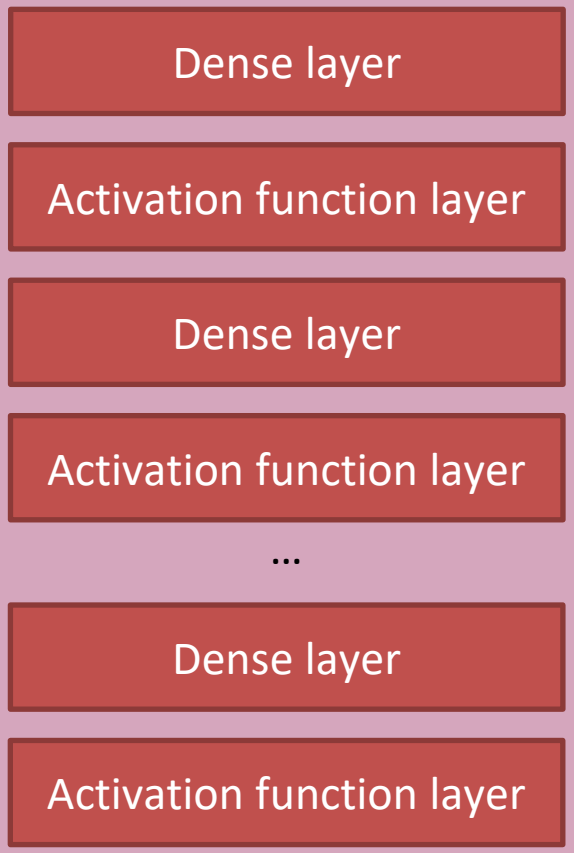

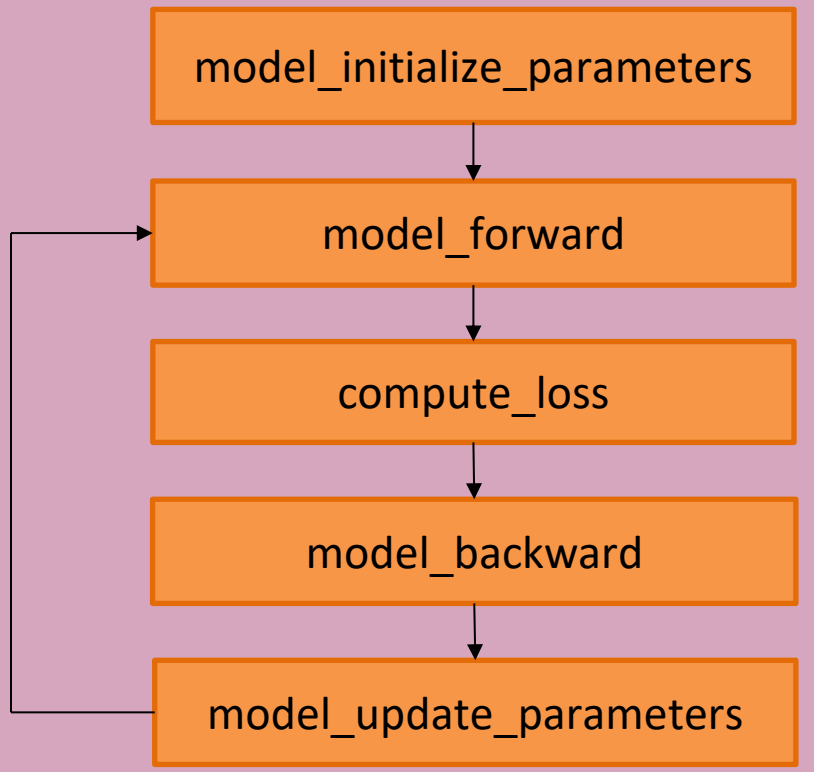

Basic Implementation (65%)

Section 1: Neural network implementation (30%)

Part 1: Linear layer (10%)

- Step 1: Linear Initialize parameters (0%)
- Step 2: Linear forward (4%)
- Step 3: Linear backward (4%)
- Step 4: Linear update parameters (2%)

Part 2: Activation function layer (10%)

- Step 1: Activation forward (5%)
- Step 2: Activation backward (5%)

Part 3: Build model (10%)

- Step 1: Model Initialize parameters (0%)
- Step 2: Model forward (4%)
- Step 3: Model backward (4%)
- Step 4: Model update parameters (2%)

Basic Implementation (65%)

Section 2: Loss function (10%)

Part 1: Binary cross-entropy loss (5%)

Part 2: Categorical cross-entropy loss (5%)

Part 3: Mean square error (0%)

Section 3: Training and prediction (25%)

Part 1: Training function & batch function (5%)

Part 2: Regression (10%)

- Baseline : MAE <= 0.05 (5%)
- Submit "**Lab4_basic_regression.gif**" (5%)

Part 3: Binary classification (10%)

- Baseline : Public f1 score >= 0.8 (5%)
- Baseline : Private f1 score >= 0.8 (5%)

Advanced Implementation (30%)

Multi-class classification

- Baseline : Public f1 score >= 0.6 (5%)
- Baseline : Private f1 score >= 0.6 (10%)
- Private Ranking (15%)

Loss function and Activation function

Warning: only the following 3 combinations are allowed!

1. Regression : linear + mse
2. Binary classification: sigmoid + cross_entropy
3. Multi-class classification: softmax + cross_entropy

Data (Simulation data)

Regression: Math function approximation

The target function to approximate is:

$$y = \sin(2 * \sin(2 * \sin(2 * \sin(x))))$$

where x is in the range [0.01, 11]

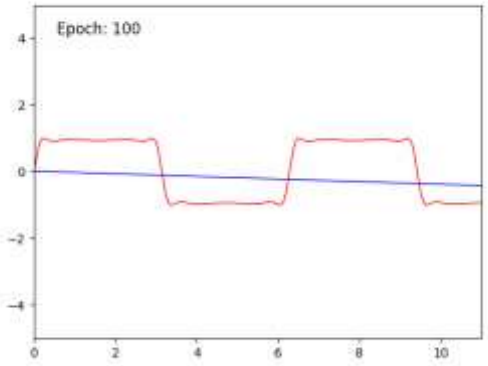

Data (OCT scans)

Binary classification: OCT scan images of retina

The dataset consists of 28x28 pixels grayscale OCT scan images of the retina, focusing on two classes: CNV (Choroidal Neovascularization) and Normal.

Details of the dataset:

- shape of x_train: (20000, 28, 28)
- shape of y_train: (20000, 1)
- shape of x_test: (5000, 28, 28)

Classes:

- CNV: label = 1
- Normal: label = 0

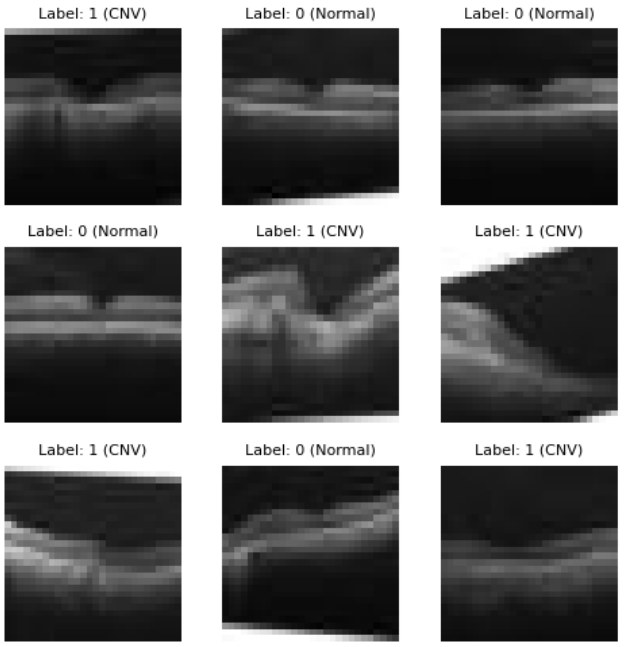

Data (OCT scans)

Multi-class classification: OCT scan images of retina

The dataset consists of 28x28 pixels grayscale OCT scan images of the retina, categorized into four classes: CNV (Choroidal Neovascularization), DME (Diabetic Macular Edema), Drusen, and Normal.

Details of the dataset:

- shape of x_train: (37754, 28, 28)
- shape of y_train: (37754,)
- shape of x_test: (3000, 28, 28)

Classes:

- CNV: label = 0
- DME: label = 1
- Drusen: label = 2
- Normal: label = 3

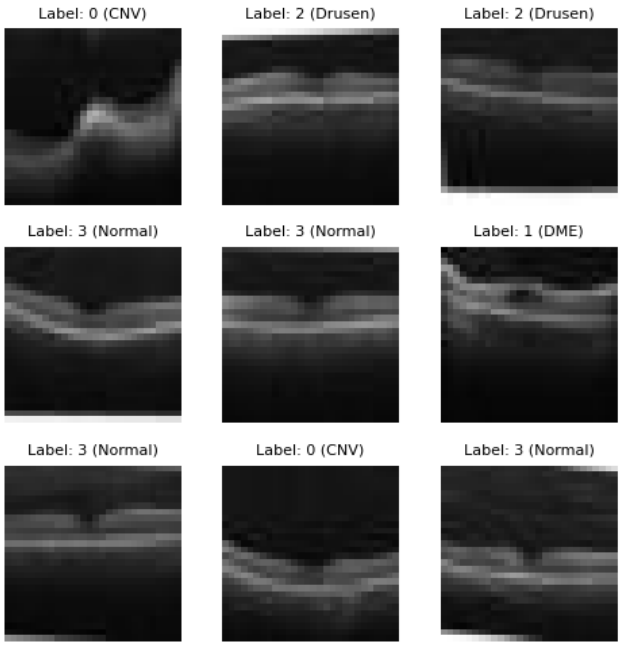

Output .csv file format

Basic Part : regression

There should be (1000+1) rows in your csv file

First row is the header [‘ID’, ’y’]

ID starts from 0, and **y** is the predicted y value

Please make sure that your output format is correct

Submit the answer (Lab4_basic_regression.csv) to Kaggle

ML2024-Lab4-BasicPart-Reg

	A	B
1	ID	y
2	0	0.389933
3	1	0.420138
4	2	0.450343
5	3	0.480549
6	4	0.510754
7	5	0.540959
8	6	0.571164
9	7	0.601369
10	8	0.631574
11	9	0.66178
12	10	0.691985

Output .csv file format

Basic Part : binary classifier

There should be (5000+1) rows in your csv file

- First row is the header [‘ID’, 'Label']
- Your prediction answer should be either 0 or 1
- **ID** starts from 0, and **Label** is the predicted answer

Please make sure that your output format is correct

Submit the answer (Lab4_basic.csv) to Kaggle

ML2024-Lab4-BasicPart-B.C.

	A	B
1	ID	Label
2	0	1
3	1	1
4	2	1
5	3	1
6	4	0
7	5	0
8	6	0

Output .csv file format

Advanced Part : multi-class classifier

There should be (3000+1) rows in your csv file

- First row is the header [‘ID’, 'Label']
- Your prediction answer should be (0~3)
- **ID** starts from 0, and **Label** is the predicted answer

Please make sure that your output format is correct

Submit the answer (Lab4_advanced.csv) to Kaggle

ML2024-Lab4-AdvancedPart

	A	B
1	ID	Label
2	0	1
3	1	3
4	2	3
5	3	1
6	4	1
7	5	0
8	6	3

Output .npy File Format

- Named as “**Lab4_output.npy**”
- This file is a dictionary that stores your output for each function.
- We will test your ”**Lab4_output.npy**” to verify the correctness of your neural networks.

```
dense_forward: <class 'tuple'>
dense_backward: <class 'tuple'>
dense_update_parameters: <class 'dict'>
sigmoid: <class 'tuple'>
relu: <class 'tuple'>
softmax: <class 'tuple'>
linear: <class 'tuple'>
sigmoid_backward: <class 'numpy.ndarray'>
relu_backward: <class 'numpy.ndarray'>
softmax_backward: <class 'numpy.ndarray'>
linear_backward: <class 'numpy.ndarray'>
model_forward_sigmoid: <class 'tuple'>
model_forward_relu: <class 'tuple'>
model_forward_softmax: <class 'tuple'>
model_backward_sigmoid: <class 'tuple'>
model_backward_relu: <class 'tuple'>
model_update_parameters: <class 'dict'>
compute_BCE_loss: <class 'numpy.float64'>
compute_CCE_loss: <class 'numpy.float64'>
```

Evaluation Metric

Regression

- MAE (mean absolute error)

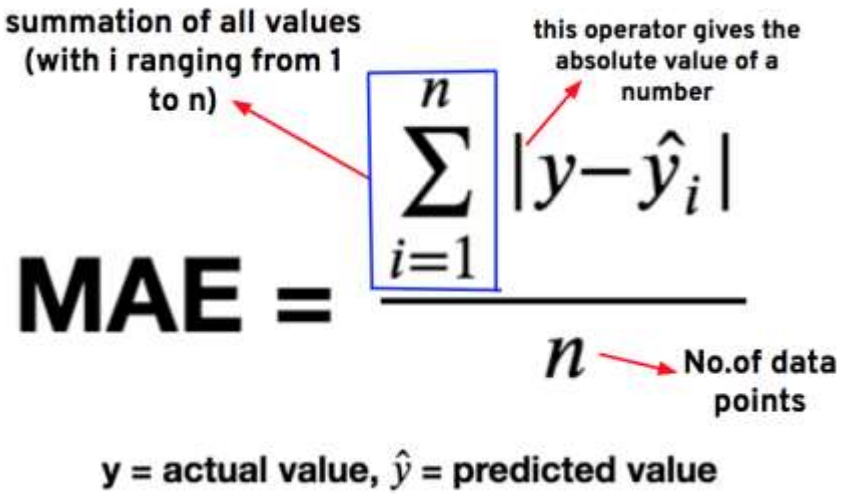

Evaluation Metric

Classification

- F1-score

$$F1\text{-}score = 2 \times \frac{(Precision \times Recall)}{(Precision + Recall)}$$

		Actual/True value	
		positive	negative
Predicted value	positive	**TP**	**FP**
	negative	**FN**	**TN**

Given Items

- Template: Lab4_template.ipynb
- Basic data (binary classifier): basic_data.npz
- Advanced data: advanced_data.npz

Template

Important notice

- Please **do not** change the code outside this code bracket in the basic part.

```
### START CODE HERE ###
...
### END CODE HERE ###
```

- Please **do not** import any other packages in both basic and advanced part
- Please **do not** change the random seed **np.random.seed(1)**.

Remember to save the code file to **Lab4.ipynb**

Kaggle

We've created competitions for 3 tasks respectively.

- Basic regression link:
  https://www.kaggle.com/competitions/ml-2024-lab-4-basic-part-fa
- Basic binary classification link:
  https://www.kaggle.com/competitions/ml-2024-lab-4-basic-part-bc
- Advanced link:
  https://www.kaggle.com/competitions/ml-2024-lab-4-advanced-part

- For regression, only **public** data are provided.
- For both binary and multi-class classification tasks, we split the testing data randomly into **public** (50%) and **private** (50%) parts, maintaining the same class distribution ratio.
- Only the public score will be visible on Kaggle.

Kaggle

- Please register your account.
- Click the ‘Join competition’ button to join.

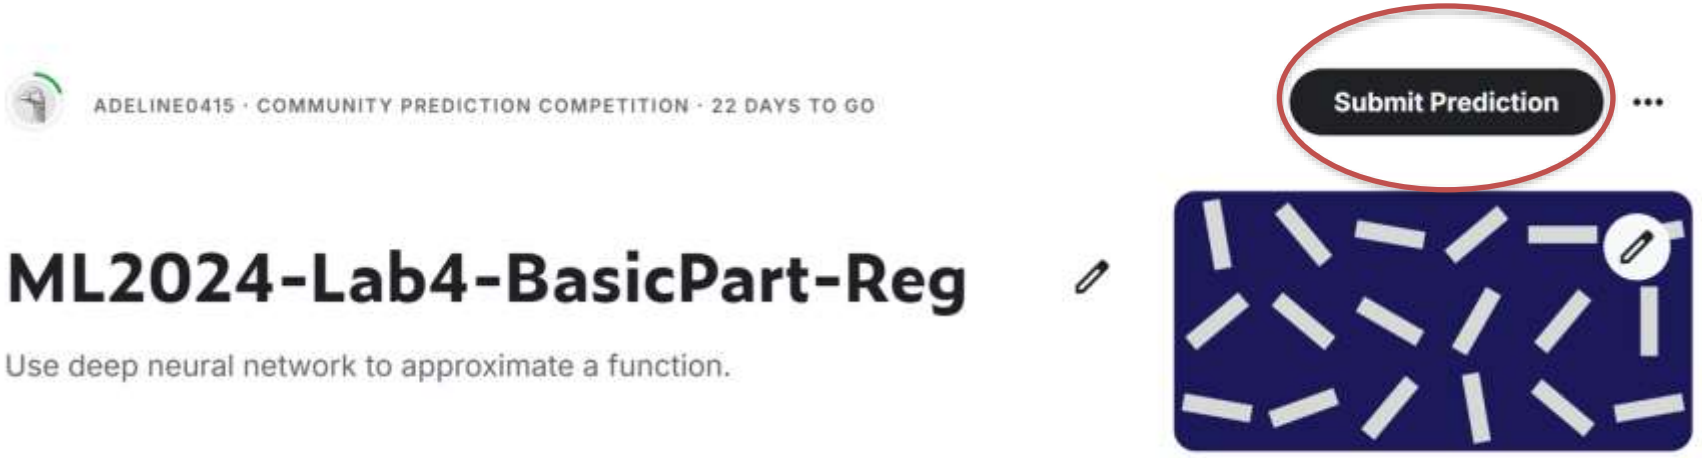

Kaggle

- After joining the competition, you should change your team name (each student is a team) to your **student ID**.
- **Please remember to SAVE CHANGES**
- You can submit 50 times per day.

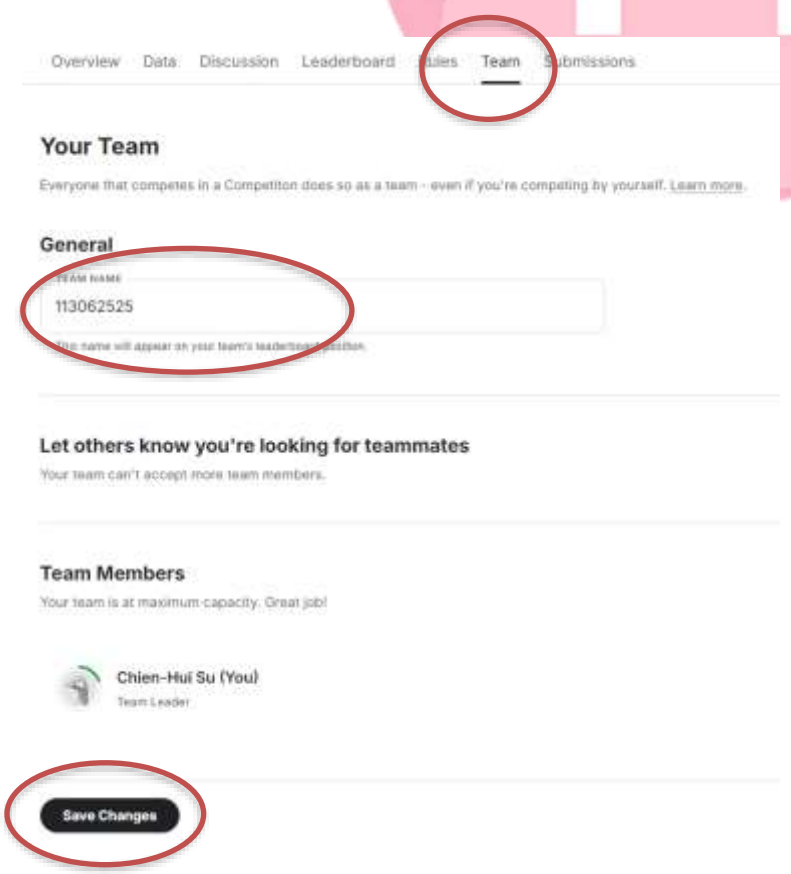

Notes: Please verify your team name on the leaderboard - changing profile name does not change team name.

Kaggle

You can manually select up to 1 submission that will count towards your final leaderboard score. If no submission is selected, Kaggle will automatically select your submission with the best public score.

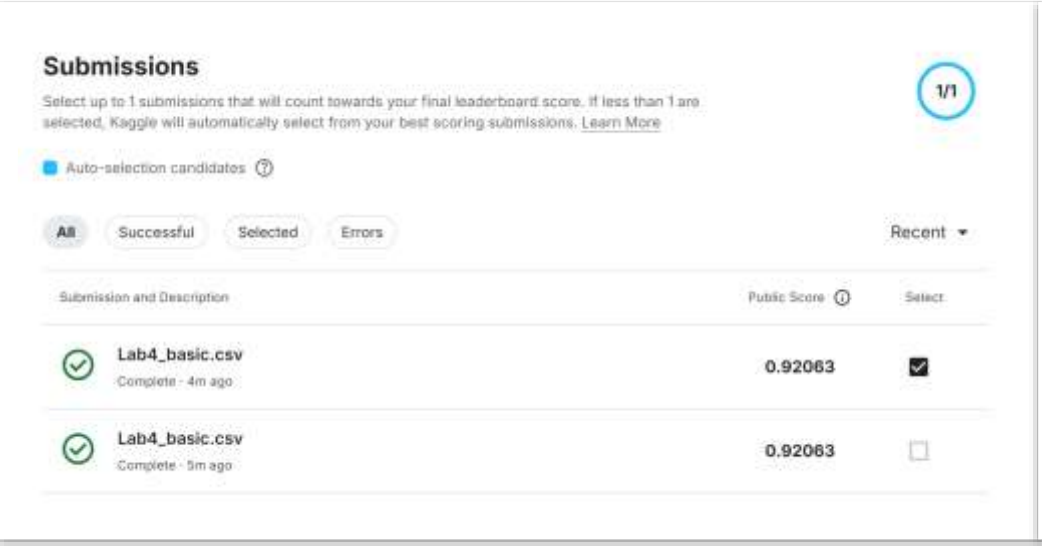

Manual-selection

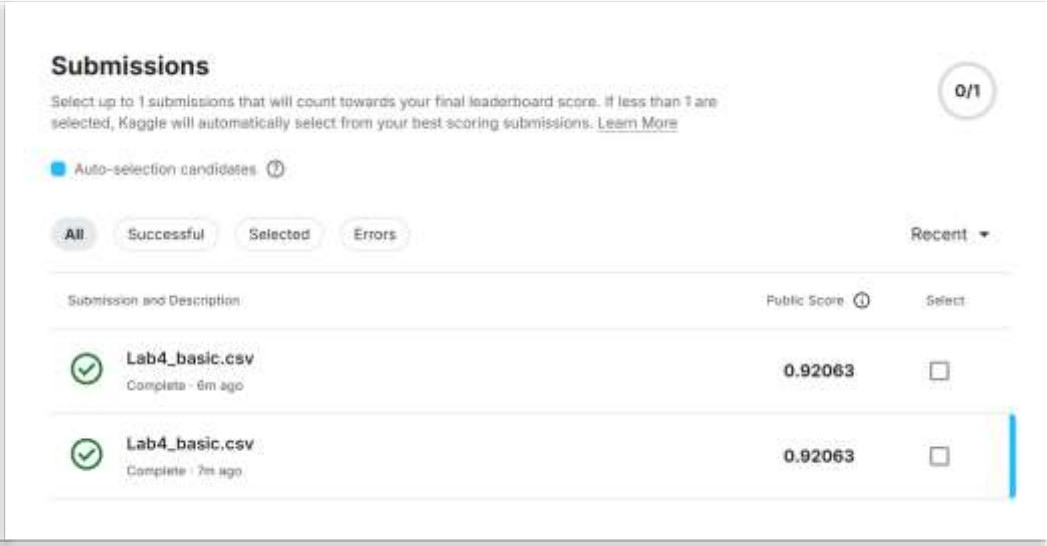

Auto-selection

Basic & Advanced Report (5%)

1. What are the key differences between sigmoid and softmax activation functions, and why did we choose them for binary and multi-class classification respectively? (1%)
2. Why does the loss oscillate during model training? (list at least 2 reasons) (2%)
3. How does changing the learning rate and batch size affect model training time? (1%)
4. Put your regression results (*lab4_basic_regression.jpg*) on report. (1%)

Notes:

1. Do not exceed 1 page!
2. Name your report file as "**Lab4_report.pdf**".

Requirement

- Do it individually! Not as a team! (team is for final project)
- Announce date: 2024/10/31
- Deadline: 2024/11/12 23:59 (Late submission is not allowed!)
- Submit the answers (csv) to corresponding Kaggle competition.
  - **ML2024-Lab4-BasicPart-Reg**
  - **ML2024-Lab4-BasicPart-B.C.**
  - **ML2024-Lab4-AdvancedPart**
- Hand in following files to eeclass in the following format (Do not compressed!)
  - Lab4.ipynb
  - Lab4_report.pdf
  - Lab4_basic_regression.gif
  - Lab4_output.npy

Penalty

0 points if any of the following conditions happened

- Plagiarism
- Late submission
- Not using a template or importing any other packages
- No submission record on Kaggle (we cannot identify who you are)
- Wrong team name on Kaggle (we cannot identify who you are)
- No code("**Lab4.ipynb**") submission on eeclass
- Your submission was not generated by your code

5 Points would be deducted if your submission format is incorrect

0 Points will be given in the Basic section 1&2 if you don't submit "**Lab4_output.npy**"

Questions?

- TA: Chia-Suan Yu 余佳軒 (adeline041503@gmail.com)
- No debugging service